

# Quant II

## Lab 2: Regression

Giacomo Lemoli

February 7, 2022

# Today's plan


- DAGs
- Regression
- Heterogeneity and effective samples
- Running regressions in practice

# Conditioning in DAGs

Conditioning on some variable  $w$  in a DAG is equivalent to do the following steps:

- If  $w$  is a collider, link all pairs of parents of  $w$  by drawing an undirected edge between them
- For any ancestor of  $w$ , if this ancestor is itself a collider, link all pairs of parents of this ancestor with undirected edges to connote induced dependencies
- Erase  $w$  from the graph and all the edges connected with  $w$


An example from the Mixtape



# DAG simulations

- Let's see these properties at work, with some simulations
- Examples and code from The Mixtape, (pp. 108-113)

Case 1: the effect of gender discrimination on women income



# DAG simulations

```
library(tidyverse)
library(stargazer)

# Set seed
set.seed(123)

# Simulate our data
tb <- tibble(
  female = ifelse(runif(10000) >= 0.5, 1, 0),
  ability = rnorm(10000),
  discrimination = female,
  occupation = 1 + 2*ability + 0*female - 2*discrimination + rnorm(10000),
  wage = 1-1*discrimination + 1*occupation + 2*ability + rnorm(10000)
)

# Estimate regressions
lm_1 <- lm(wage ~ female, tb)
lm_2 <- lm(wage ~ female + occupation, tb)
lm_3 <- lm(wage ~ female + occupation + ability, tb)
```

# DAG simulations


# Compare

```
stargazer(lm_1, lm_2, lm_3, type = "text",
          column.labels = c("Biased unconditional",
                           "Biased",
                           "Unbiased Conditional"))
```

```
##
## =====
##                               Dependent variable:
## -----
##                               wage
##                               Biased
##                               Unbiased Conditional
##                               (1)          (2)          (3)
## -----
## female                -3.066***          0.587***          -1.050***
##                        (0.085)          (0.030)          (0.028)
##
## occupation                1.796***          0.987***
##                        (0.006)          (0.010)
##
## ability                                2.033***
##                                (0.022)
##
## Constant                2.023***          0.222***          1.025***
##                        (0.060)          (0.020)          (0.017)
## -----
## Observations                10,000          10,000          10,000
## R2                        0.114          0.912          0.952
## Adjusted R2                0.114          0.912          0.952
## Residual Std. Error      4.265 (df = 9998)      1.347 (df = 9997)      0.994 (df = 9996)
## F Statistic      1,292.306*** (df = 1; 9998) 51,551.530*** (df = 2; 9997) 65,927.470*** (df = 3; 9996)
## =====
## Note: *p<0.1; **p<0.05; ***p<0.01
```

# DAG simulations

Case 2: Talent and beauty





# DAG simulations


```
library(tidyverse)

# Set seed
set.seed(3444)

# Simulate data
star_is_born <- tibble(
  beauty = rnorm(2500),
  talent = rnorm(2500),
  score = beauty + talent,
  c85 = quantile(score, .85),
  star = ifelse(score >= c85, 1, 0)
)
```


# DAG simulations

```
star_is_born %>%  
  lm(beauty ~ talent, .) %>%  
  ggplot(aes(x = talent, y = beauty)) +  
  geom_point(size = 0.5, shape = 23) + xlim(-4, 4) + ylim(-4, 4) +  
  theme_bw()
```



# DAG simulations

```
star_is_born %>%  
  filter(star == 1) %>% lm(beauty ~ talent, .) %>%  
  ggplot(aes(x = talent, y = beauty)) +  
  geom_point(size = 0.5, shape = 23) + xlim(-4, 4) + ylim(-4, 4) +  
  theme_bw()
```



# Conclusion

- Don't control for/condition on colliders
- Endogenous sample selection is a form of collider bias!
  - See discussion in Knox et al (2020) on admin data

Additional resources:

- Elwert and Winship (2014), Endogenous selection bias: the problem of conditioning on a collider variable
- Knox, Lucas, and Cho (2022), Testing causal theories with learned proxies
- Schneider (2020), Collider bias in economic history research

# Frisch-Waugh-Lovell theorem: a refresher

- We have a linear model with  $K$  covariates. In matrix form:  $y = X'\beta + \epsilon$
- FWL gives a formula for the OLS estimate of the  $k^{th}$  coefficient.

$$\hat{\beta}_k = (X'_k M_{[X_{-k}]} X_k)^{-1} X'_k M_{[X_{-k}]} y$$

- In other words we can do the following:
  - Regress the individual variable  $X_k$  on all the other covariates and take the residuals
  - Regress the outcome variable  $y$  on all the covariates, except  $X_k$ , and take the residuals
  - Regress the residuals of  $y$  on the residuals for  $X$
- Note that to get  $\hat{\beta}_k$  it is enough to regress the non-residualized  $y$  on residualized  $X_k$  (because the matrix  $M$  is idempotent), but the SE won't be right
- Useful because typically we are interested in just one regressor (e.g. a treatment indicator), so we can reduce the dimensionality of the model

# FWL in practice

```
# Import a dataset
data("mtcars")

# Multivariate regression
fit <- lm(mpg ~ cyl + drat + wt, mtcars)

# FWL
resy <- lm(mpg ~ drat + wt, mtcars) %>% residuals()
resx <- lm(cyl ~ drat + wt, mtcars) %>% residuals()
fit2 <- lm(resy ~ resx)


# Compare results
out <- c(coefficients(fit)["cyl"], coefficients(fit2)["resx"])
names(out) <- c("Multivariate", "Univariate Residualized")
out

##           Multivariate Univariate Residualized
##           -1.509577      -1.509577
```

# FWL in practice

With residual-residual plots we can show the relationship between two variables while “controlling” for the others

```
as.data.frame(cbind(resy, resx)) %>% rename(mpg_res = resy, cyl_res = resx) %>%  
  ggplot(aes(x=cyl_res, y=mpg_res)) + geom_point(size=1.5, colour="black", shape=21) +  
  geom_smooth(method="lm") +  
  labs(x = "Cyl, res.", y = "MPG, res.") + theme_bw()
```



# Heterogeneous treatment effects

With heterogeneous treatment effects, OLS estimators have some features to be aware of

- Two groups in the population, A and B. Represented by covariate  $X_i \in \{A, B\}$ . We draw a random sample. We want to estimate the ATT.
- There is treatment effect heterogeneity by group:  $ATT_A \neq ATT_B$ . CIA holds:  $(Y_{0i}, Y_{1i}) \perp D_i | X_i$
- CIA intuitively means we can do the following:
  - 1 Fix a value of  $X_i$
  - 2 Estimate the effect of  $D$  on  $Y$  for units with that value of  $X_i$ . This is a causal estimate
  - 3 Repeat for all values of  $X_i$
  - 4 Aggregate all these causal estimates: a weighted average



# Heterogeneous treatment effects

Call  $\alpha_A$  the share of units in group A among the treated in the population, and  $\alpha_B$  the share of units in group B. Then,

$$ATT = \alpha_A * ATT_A + \alpha_B * ATT_B$$

$$\alpha_A + \alpha_B = 1$$

We can rewrite as:

$$ATT = \alpha_A * ATT_A + \alpha_B * ATT_B$$

$$ATT = P(X_i = A | D_i = 1) \{E[Y_{1i} - Y_{0i} | X_i = A, D_i = 1]\} + \\ P(X_i = B | D_i = 1) \{E[Y_{1i} - Y_{0i} | X_i = B, D_i = 1]\}$$

The terms in red are the weights we use in the weighted average: the share of units with the same ATT.

By CIA, the terms in black are identified. So we can estimate ATT with sample analogues.

# Heterogeneous treatment effects

An unbiased estimator is

$$A\hat{T}T = \frac{\sum_{x \in \{A, B\}} \hat{\delta}_x \hat{P}(D_i = 1 | X_i = x) \hat{P}(X_i = x)}{\sum_{x \in \{A, B\}} \hat{P}(D_i = 1 | X_i = x) \hat{P}(X_i = x)}$$

OLS regression also estimates a weighted average of individual-unit treatment effects (proof in lecture slides and MHE)

$$A\hat{T}T_R = \frac{\sum_{x \in \{A, B\}} \hat{\delta}_x \hat{P}(D_i = 1 | X_i = x) (1 - \hat{P}(D_i = 1 | X_i = x)) \hat{P}(X_i = x)}{\sum_{x \in \{A, B\}} \hat{P}(D_i = 1 | X_i = x) (1 - \hat{P}(D_i = 1 | X_i = x)) \hat{P}(X_i = x)}$$

This is not the weighted average we have started with: instead of weighting more the group that represents more units, it weights more the group where the treatment status has higher variance.

How distributed the treatment needs to be in a group in order for it to have the highest weight?

# Effective sample

- Aronow and Samii (2016) show that under some assumptions about the functional form of the treatment assignment the following result holds:

$$\hat{\beta} \xrightarrow{P} \frac{E[w_i \tau_i]}{E[w_i]}, \text{ where } w_i = (D_i - E[D_i|X_i])^2$$

so that

$$E[w_i|X_i] = E[D_i - E[D_i|X_i]|X_i]^2 = \text{Var}[D_i|X_i]$$

- This result implies that regression re-weights units in ways that are not detectable at first sight
- Units in groups/Covariates strata where the treatment has a higher conditional variance receive more weight
- Equivalent to run the regression on an *effective* sample different from the one we think we are working with
- To characterize the effective sample we can estimate the  $w_i$ s

# Effective sample

Let's rewrite the last expression.

$$E[w_i|X_i] = E[(D_i - E[D_i|X_i]|X_i)^2] = \text{Var}[D_i|X_i]$$

- If we assume linearity of the treatment assignment in  $X_i$ , the weight is equal to the square of the residual of regressing the treatment indicator on  $X_i$
- Intuition: higher conditional variance of treatment  $\implies$  treatment has more residual variance not explained by the covariates  $\implies$  higher error term
- The regression exploits as much as possible this identifying variation
- We can estimate the regression weights by the following procedure:
  - 1 Run the regression  $D_i = X_i\gamma + e_i$
  - 2 Take residual  $\hat{e}_i = D_i - X_i\hat{\gamma}$  and square it

# Effective sample: example

- Let's study the effective sample in an actual paper
- Egan and Mullin (2012) look at how people form their attitudes based on personal experiences
- They use local weather variation to estimate the effect of experiencing weather changes on beliefs about global warming
- To understand the effective sample, we need to ask where weather is most variable (conditional on covariates)

# Effective sample: example

```
# Import the data
library(haven)
d <- read_dta("gwdataset.dta")

# Import state IDs
zips <- read_dta("zipcodetostate.dta")
zips <- zips %>% select(c(statenum, statefromzipfile)) %>% unique()
zips <- zips %>% filter(!(statenum == 8 & statefromzipfile == "NY"))

# Import population data
pops <- read.csv("population_ests_2013.csv")

# Format
pops$state <- tolower(pops$NAME)
d$getwarmord <- as.double(d$getwarmord)
```

# Effective sample: example

```
# Estimate primary model of interest:
```

```
d$doi <- factor(d$doi)
```

```
d$statenum <- factor(d$statenum)
```

```
d$wbnid_num <- factor(d$wbnid_num)
```

```
Y <- "getwarmord"
```

```
D <- "ddt_week"
```

```
X <- names(d)[c(15,17,42:72)]
```

```
reg_formula <- paste0(Y, "~", D, "+", paste0(X, collapse = "+"))
```

```
reg_out <- lm(as.formula(reg_formula), d)
```

```
# Or
```

```
out <- lm(getwarmord~ddt_week+educ_hsless+educ_coll+educ_postgrad+  
  educ_dk+party_rep+party_leanrep+party_leandem+  
  party_dem+male+raceeth_black+raceeth_hisp+  
  raceeth_notwbh+raceeth_dkref+age_1824+age_2534+  
  age_3544+age_5564+age_65plus+age_dk+ideo_vcons+  
  ideo_conservative+ideo_liberal+ideo_vlib+ideo_dk+  
  attend_1+attend_2+attend_3+attend_5+attend_6+  
  attend_9+as.factor(doi)+as.factor(statenum)+  
  as.factor(wbnid_num),d)
```

# Base Model

```
summary(reg_out)$coefficients[1:10,]
```

##	Estimate	Std. Error	t value	Pr(> t )
## (Intercept)	1.945740062	0.771478843	2.5220913	0.01169077
## ddt_week	0.004857915	0.002475887	1.9620908	0.04979656
## wbnid_num3103	0.843451519	0.922666490	0.9141456	0.36067588
## wbnid_num3154	1.575071541	0.973391215	1.6181280	0.10568587
## wbnid_num3159	1.903629413	1.021302199	1.8639237	0.06237963
## wbnid_num3804	1.406498119	0.794035963	1.7713280	0.07655528
## wbnid_num3810	1.330878449	0.806312016	1.6505750	0.09887602
## wbnid_num3811	1.082204367	0.798796489	1.3547936	0.17553267
## wbnid_num3812	1.219327925	0.803974284	1.5166255	0.12941222
## wbnid_num3813	0.986084952	0.829563706	1.1886790	0.23461152



# Estimate the weights

```
# Regress treatment indicator on the vector of covariates  
D_formula <- paste0(D, "~", paste0(X, collapse = "+"))  
outD <- lm(as.formula(D_formula),d)  
  
# Extract the residuals and take their square  
eD2 <- residuals(outD)^2
```

# Effective sample statistics

- We can use these estimated weights to study the effective sample

```
# Take some relevant variables
```

```
compare_samples<- d[, c("wave", "ddt_week", "ddt_twoweeks",  
  "ddt_threeweeks", "party_rep", "attend_1", "ideo_conservative",  
  "age_1824", "educ_hsless")]
```

```
# Compute statistics with and without weights
```

```
compare_samples <- t(apply(compare_samples,2,function(x)  
  c(mean(x),sd(x),weighted.mean(x,eD2),  
    sqrt(weighted.mean((x-weighted.mean(x,eD2))^2,eD2))))  
colnames(compare_samples) <- c("Nominal Mean", "Nominal SD",  
  "Effective Mean", "Effective SD")
```

# Effective Sample Statistics

```
compare_samples
```

##	Nominal Mean	Nominal SD	Effective Mean	Effective SD
## wave	3.09693726	1.4252527	3.20788200	1.5609143
## ddt_week	3.83548593	5.9047249	5.11579140	10.8980228
## ddt_twoweeks	3.85505617	5.4572382	5.00137435	9.2262827
## ddt_threeweeks	3.96719696	4.7689594	5.10859485	8.4348180
## party_rep	0.29527208	0.4561989	0.28978321	0.4536617
## attend_1	0.11433244	0.3182383	0.12343459	0.3289354
## ideo_conservative	0.31132917	0.4630715	0.29325249	0.4552532
## age_1824	0.07195956	0.2584402	0.06881146	0.2531333
## educ_hsless	0.34151056	0.4742516	0.31219962	0.4633908

# Effective sample maps

- We can depict the samples visually
- What places in the US are “over-represented” in the effective samples?

# Effective sample maps

```
# Construct the "effective sample weights" for each state
wts_by_state <- tapply(eD2, d$statenum, sum)
wts_by_state <- wts_by_state/sum(wts_by_state)*100
wts_by_state <- data.frame(eff = wts_by_state,
                           statenum = as.numeric(names(wts_by_state)))

# Merge to the state name variable
data_for_map <- merge(wts_by_state, zips, by="statenum")

# Construct the "nominal sample weights" for each state
wts_by_state <- tapply(rep(1,6726),d$statenum,sum)
wts_by_state <- wts_by_state/sum(wts_by_state)*100
wts_by_state <- data.frame(nom = wts_by_state,
                           statenum = as.numeric(names(wts_by_state)))

# Add to the other data
data_for_map <- merge(data_for_map, wts_by_state, by="statenum")
```

# Effective sample maps

```
# Get correct state names
require(maps,quietly=TRUE)
data(state.fips)

# Add them to the dataset
data_for_map <- left_join(data_for_map, state.fips,
                          by = c("statefromzipfile" = "abb"))

# More data prep
data_for_map$state <- sapply(as.character(data_for_map$polynome),
                             function(x)strsplit(x,":")[[1]][1])
data_for_map <- data_for_map %>% group_by(statefromzipfile) %>%
  summarise_all(first) %>% ungroup() %>% select(-polynome)

# Diff between nominal and effective weights
data_for_map$diff <- data_for_map$eff - data_for_map$nom

# Merge with population data
data_for_map <- left_join(data_for_map, pops, by="state")

# Actual "weight" of each state in the US
data_for_map$pop_pct <- data_for_map$POPESTIMATE2013/sum(
  data_for_map$POPESTIMATE2013)*100

# Different representativity of the two samples
data_for_map <- mutate(data_for_map,
                       pop_diff_eff = eff - pop_pct,
                       pop_diff_nom = nom - pop_pct)
data_for_map <- mutate(data_for_map,
                       pop_diff = pop_diff_eff - pop_diff_nom)

require(ggplot2,quietly=TRUE)
state_map <- map_data("state")
```

# More setup


```
# Plot the weights in each sample
plot_eff <- ggplot(data_for_map, aes(map_id = state)) +
  geom_map(aes(fill=eff), map = state_map) +
  expand_limits(x= state_map$long, y = state_map$lat) +
  scale_fill_continuous("% Weight", limits=c(0,17), low="white", high="black") +
  labs(title = "Effective Sample") +
  theme(legend.position=c(.2,.1), legend.direction = "horizontal",
        axis.line = element_blank(), axis.text = element_blank(),
        axis.ticks = element_blank(), axis.title = element_blank(),
        panel.background = element_blank(),
        plot.background = element_blank(),
        panel.border = element_blank(),
        panel.grid = element_blank())

plot_nom <- ggplot(data_for_map, aes(map_id = state)) +
  geom_map(aes(fill=nom), map = state_map) +
  expand_limits(x=state_map$long, y=state_map$lat) +
  scale_fill_continuous("% Weight", limits=c(0,17), low="white", high="black") +
  labs(title="Nominal Sample") +
  theme(legend.position=c(.2,.1), legend.direction = "horizontal",
        axis.line = element_blank(), axis.text = element_blank(),
        axis.ticks = element_blank(), axis.title = element_blank(),
        panel.background = element_blank(),
        plot.background = element_blank(),
        panel.border = element_blank(), panel.grid = element_blank())
```


# Maps

```
require(gridExtra,quietly=TRUE)  
grid.arrange(plot_nom,plot_eff,ncol=2)
```

Nominal Sample



Effective Sample






# Setup comparison plot

```
plot_diff <- ggplot(data_for_map,aes(map_id=state)) +  
  geom_map(aes(fill=diff), map = state_map) +  
  expand_limits(x = state_map$long, y = state_map$lat) +  
  scale_fill_gradient2("% Weight", low = "red", mid = "white", high = "black") +  
  labs(title = "Effective Weight minus Nominal Weight") +  
  theme(legend.position=c(.2,.1),legend.direction = "horizontal",  
        axis.line = element_blank(), axis.text = element_blank(),  
        axis.ticks = element_blank(), axis.title = element_blank(),  
        panel.background = element_blank(),  
        plot.background = element_blank(),  
        panel.border = element_blank(), panel.grid = element_blank())
```

## Difference in weights

plot\_diff

## Effective Weight minus Nominal Weight



# Linear regression in practice

- In R the command `lm()` implements the OLS estimation
- Now a variety of packages handle robust SEs and FEs much better (more on this in the coming weeks)
  - `estimatr::lm_robust()`. Good for design of randomized experiments
  - `fixest::feols()`. Good for panel data with many FEs
  - `lfe::felm()`. Once popular for panel data, now on the way to be abandoned
- Stata: `reg[ress]`. Better default than `lm()` but also superseded in contemporary empirical research
  - `xtreg`, `areg`, `reghdfe` all used in empirical research due to the handling of FEs, the latter seems to be taking over
- Packages to export output in tables: `stargazer`, `modelsummary`
- In Stata: `outreg2`, `estout`